

Breast Cancer Tumor Detection Using MRI Dataset



Email: - 2k22.cse.2213298@gmail.com

Name: - Aryan Katiyar

[Infosys Springboard 5.0]

CODE → // TASKS PERFORMED :-

//1. COUNT OF BENIGN AND MALIGNANT IMAGES FROM THE TEST DATASET

// 2. DATA AUGMENTATION PERFORMED ON TEST DATASET

// 3. RANDOM CROP AND NORMALIZATION PERFORMED ON TRAIN AND VALIDATION DATASET

```
import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from collections import Counter
import numpy as np
import torch

# Define transformations for the training set (with data augmentation)
train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224),          # Randomly crop the
    image and resize to 224x224
    transforms.RandomHorizontalFlip(),          # Randomly flip the
    images horizontally
    transforms.RandomVerticalFlip(),            # Randomly flip the
    images vertically
    transforms.RandomRotation(20),              # Rotate the images by
    up to 20 degrees
    transforms.ColorJitter(brightness=0.2, contrast=0.2,
    saturation=0.2), # Add random color jitter
    transforms.RandomAffine(degrees=0, translate=(0.1, 0.1)), #
    Randomly translate the image
    transforms.RandomPerspective(distortion_scale=0.5, p=0.5), #
    Random perspective transformation
    transforms.ToTensor(),                      # Convert the image to a
    PyTorch tensor
    transforms.GaussianBlur(kernel_size=(5, 9), sigma=(0.1, 5)), #
    Apply Gaussian blur with random kernel
    transforms.RandomErasing(p=0.5, scale=(0.02, 0.2)),          #
    Randomly erase a portion of the image (applied to tensor)
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
    0.225]) # Normalize based on ImageNet
])
```

```

# Define transformations for test and validation sets (only resize and
normalization)
test_val_transform = transforms.Compose([
    transforms.Resize((224, 224)),      # Resize the image to 224x224
    transforms.ToTensor(),              # Convert the image to a
PyTorch tensor
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225]) # Normalize based on ImageNet
])

# Path to subfolders
data_dir_train = '/content/drive/MyDrive/Tumor_Trace/clasificacion-
roi/train'
data_dir_test = '/content/drive/MyDrive/Tumor_Trace/clasificacion-
roi/test'
data_dir_val = '/content/drive/MyDrive/Tumor_Trace/clasificacion-
roi/val'

# Load the datasets using ImageFolder
train_dataset = datasets.ImageFolder(root=data_dir_train,
transform=train_transform)
test_dataset = datasets.ImageFolder(root=data_dir_test,
transform=test_val_transform)
val_dataset = datasets.ImageFolder(root=data_dir_val,
transform=test_val_transform)

# Count the total number of images in each class for the training
dataset
class_counts = Counter(train_dataset.targets)

benign_count = class_counts[train_dataset.class_to_idx['Benign']]
malignant_count = class_counts[train_dataset.class_to_idx['Malignant']]

# Print the counts
print(f"Total Benign Images: {benign_count}")
print(f"Total Malignant Images: {malignant_count}")

# Plotting the graph between total number of benign and malignant
images
plt.figure(figsize=(6, 4))
plt.bar(['benign', 'malignant'], [benign_count, malignant_count],
color=['blue', 'red'])
plt.title('Distribution of Benign and Malignant Images')
plt.xlabel('Classes')
plt.ylabel('Number of Images')
plt.show()

```

```

# Visualizing some augmented images from the training dataset
def imshow(img):
    img = img / 2 + 0.5 # unnormalize the image (after normalization
                        in transform)
    np_img = img.numpy()
    plt.imshow(np.transpose(np_img, (1, 2, 0))) # convert from tensor
to image
    plt.show()

# Create DataLoader for augmented image visualization (from training
dataset)
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

# Get a batch of images from the train dataloader
data_iter = iter(train_loader)
images, labels = next(data_iter)

# Show a few augmented images from training set
print("Augmented Images from Training Set:")
for i in range(5): # Display first 5 images from the batch
    imshow(images[i])


# Create DataLoader for test and validation datasets (no augmentation,
only resize and normalization)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)

```


OUTPUT➔

Total Benign Images: 5559


Total Malignant Images: 14895




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



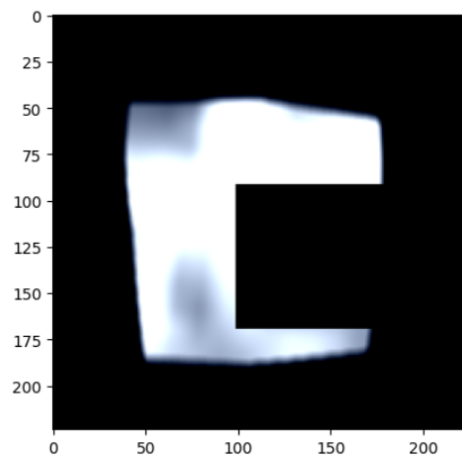
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Augmented Images from Training Set:




WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



CODE → frequency Vs input value

```
import os
import matplotlib.pyplot as plt
import numpy as np
from PIL import Image

# Define the path to the specific image
image_path = '/content/drive/MyDrive/Tumor_Trace/clasification-roi/test/Benign/BreaDM-Be-1810/SUB1/p-030.jpg' # Ensure this path is correct

# Load the image
image = Image.open(image_path).convert('L') # Ensure grayscale conversion


# Convert the image to a NumPy array
image_array = np.array(image)

# Flatten the image into a 1D array for easier plotting
flat_image_array = image_array.flatten()

# Plotting the histogram of pixel values
plt.figure(figsize=(8, 6))
plt.hist(flat_image_array, bins=256, range=(0, 255), color='blue', alpha=0.7)
plt.title(f'Pixel Value Distribution for Image: {os.path.basename(image_path)}')
plt.xlabel('Pixel Value')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()

# Print the most frequent pixel values
unique, counts = np.unique(flat_image_array, return_counts=True)
frequencies = np.asarray((unique, counts)).T
print(frequencies)
```

OUTPUT→



55 0 03

HOG IMPLEMENTATION PERFORMED

Code:-

```
import matplotlib.pyplot as plt

from skimage.feature import hog

from skimage.io import imread

from skimage.color import rgb2gray

from skimage.transform import resize

# Read image

image = imread('/kaggle/input/bctumor/test/Malignant/BreaDM-Ma-1916/SUB7/p-039.jpg')

# Resize image to 224x224

image_resized = resize(image, (224, 224))

# Convert image to grayscale if it has multiple channels (e.g., RGB)

if len(image_resized.shape) > 2:

    image_resized = rgb2gray(image_resized)

# Compute HOG features and HOG image

hog_features, hog_image = hog(image_resized,

                               orientations=9,

                               pixels_per_cell=(8, 8),

                               cells_per_block=(2, 2),
```

```
        block_norm='L2-Hys',

        visualize=True)

# Visualize the original and HOG images

fig, ax = plt.subplots(1, 2, figsize=(12, 6), sharex=True, sharey=True)

ax[0].imshow(image_resized, cmap='gray')


ax[0].set_title('Resized Original Image (224x224)')

ax[1].imshow(hog_image, cmap='gray')

ax[1].set_title('HOG Image')

plt.show()
```

Output:-



Convolution Performed:-

Code:-

```
import cv2

import numpy as np

import matplotlib.pyplot as plt

def calculate_lbp(gray_image):
    """
    Function to calculate the Local Binary Pattern (LBP) of a grayscale
    image.

    == Input ==

    gray_image: A grayscale image as a 2D NumPy array

    == Output ==

    img_lbp: LBP image of the same size
    """

    # Initialize the output LBP image with zeros

    img_lbp = np.zeros_like(gray_image)

    # Define the 3x3 neighborhood size

    neighboor = 3

    # Loop through each pixel (excluding the border pixels)
```

```

for ih in range(1, gray_image.shape[0] - 1):

    for iw in range(1, gray_image.shape[1] - 1):

        # Extract the 3x3 block of pixels centered at (ih, iw)

        block = gray_image[ih-1:ih+2, iw-1:iw+2]

        # Get the center pixel value

        center = block[1, 1]

        # Perform binary comparison and compute LBP value

        binary_pattern = (block >= center).astype(np.uint8)

        binary_pattern[1, 1] = 0 # Ignore the center pixel itself

        # Flatten and convert the binary pattern to a decimal value

        lbp_value = binary_pattern.flatten()

        lbp_value = np.delete(lbp_value, 4) # Remove center pixel
value

        lbp_decimal = np.sum(lbp_value * (2**np.arange(8)))

        # Store the LBP value in the output image

        img_lbp[ih, iw] = lbp_decimal

    return img_lbp

# Load a sample image and convert it to grayscale

```

```
image = cv2.imread('/kaggle/input/bctumor/test/Malignant/BreaDM-Ma-1916/SUB7/p-038.jpg') # Replace with the path to your image

gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Resize the image to 224x224 pixels

image = cv2.resize(image, (224, 224))

# Calculate the LBP of the image

lbp_image = calculate_lbp(gray_image)

# Display the original and LBP images

plt.figure(figsize=(10, 5))

# Original color image

plt.subplot(1, 2, 1)

plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB)) # Convert BGR to RGB for correct color display

plt.title('Original Image')

# LBP image (grayscale)


plt.subplot(1, 2, 2)

plt.imshow(lbp_image, cmap='gray')

plt.title('LBP Image')

plt.show()
```

Output:-



Mean,Median,Varience based LBP

is performed with histogtogram output

code:-

```
import numpy as np

import matplotlib.pyplot as plt

from PIL import Image

def mean_based_lbp(image_path):

    image = Image.open(image_path).convert('L').resize((224, 224))

    img_array = np.array(image)

    rows, cols = img_array.shape

    lbp_image = np.zeros((rows, cols), dtype=np.uint8)

    for i in range(1, rows - 1):

        for j in range(1, cols - 1):

            neighborhood = img_array[i-1:i+2, j-1:j+2]

            center_pixel = img_array[i, j]

            mean_value = (np.sum(neighborhood) - center_pixel) / 8

            surrounding_pixels = np.delete(neighborhood.flatten(), 4)

            binary_pattern = (surrounding_pixels >=
mean_value).astype(int)

            lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))).sum()
```

```

lbp_image_normalized = (lbp_image / lbp_image.max()) * 255

plt.figure(figsize=(14, 7))

plt.subplot(1, 3, 1)

plt.imshow(img_array, cmap='gray')

plt.title('Original Image')

plt.subplot(1, 3, 2)

plt.imshow(lbp_image_normalized, cmap='gray')

plt.title('Mean-based LBP Image')

plt.subplot(1, 3, 3)

plt.hist(lbp_image_normalized.ravel(), bins=256, range=(0, 256),
color='black')

plt.title('Histogram of Mean-based LBP')

plt.show()

def median_based_lbp(image_path):

    image = Image.open(image_path).convert('L').resize((224, 224))

    img_array = np.array(image)

    rows, cols = img_array.shape

    lbp_image = np.zeros((rows, cols), dtype=np.uint8)

    for i in range(1, rows - 1):

        for j in range(1, cols - 1):

            neighborhood = img_array[i-1:i+2, j-1:j+2]

            center_pixel = img_array[i, j]

```



```

        median_value = np.median(np.delete(neighborhood.flatten(),
4))

        surrounding_pixels = np.delete(neighborhood.flatten(), 4)

        binary_pattern = (surrounding_pixels >=
median_value).astype(int)

        lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))).sum()

lbp_image_normalized = (lbp_image / lbp_image.max()) * 255

plt.figure(figsize=(14, 7))

plt.subplot(1, 3, 1)

plt.imshow(img_array, cmap='gray')

plt.title('Original Image')

plt.subplot(1, 3, 2)

plt.imshow(lbp_image_normalized, cmap='gray')

plt.title('Median-based LBP Image')

plt.subplot(1, 3, 3)

plt.hist(lbp_image_normalized.ravel(), bins=256, range=(0, 256),
color='black')

plt.title('Histogram of Median-based LBP')

plt.show()

def variance_based_lbp(image_path):

    image = Image.open(image_path).convert('L').resize((224, 224))

    img_array = np.array(image)

    rows, cols = img_array.shape

```

```

lbp_image = np.zeros((rows, cols), dtype=np.uint8)

for i in range(1, rows - 1):
    for j in range(1, cols - 1):
        neighborhood = img_array[i-1:i+2, j-1:j+2]
        center_pixel = img_array[i, j]
        variance_value = np.var(np.delete(neighborhood.flatten(),
4))

        surrounding_pixels = np.delete(neighborhood.flatten(), 4)

        binary_pattern = (surrounding_pixels >=
variance_value).astype(int)

        lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))).sum()

lbp_image_normalized = (lbp_image / lbp_image.max()) * 255

plt.figure(figsize=(14, 7))

plt.subplot(1, 3, 1)

plt.imshow(img_array, cmap='gray')

plt.title('Original Image')

plt.subplot(1, 3, 2)

plt.imshow(lbp_image_normalized, cmap='gray')

plt.title('Variance-based LBP Image')

plt.subplot(1, 3, 3)

plt.hist(lbp_image_normalized.ravel(), bins=256, range=(0, 256),
color='black')

```

```
plt.title('Histogram of Variance-based LBP')

plt.show()

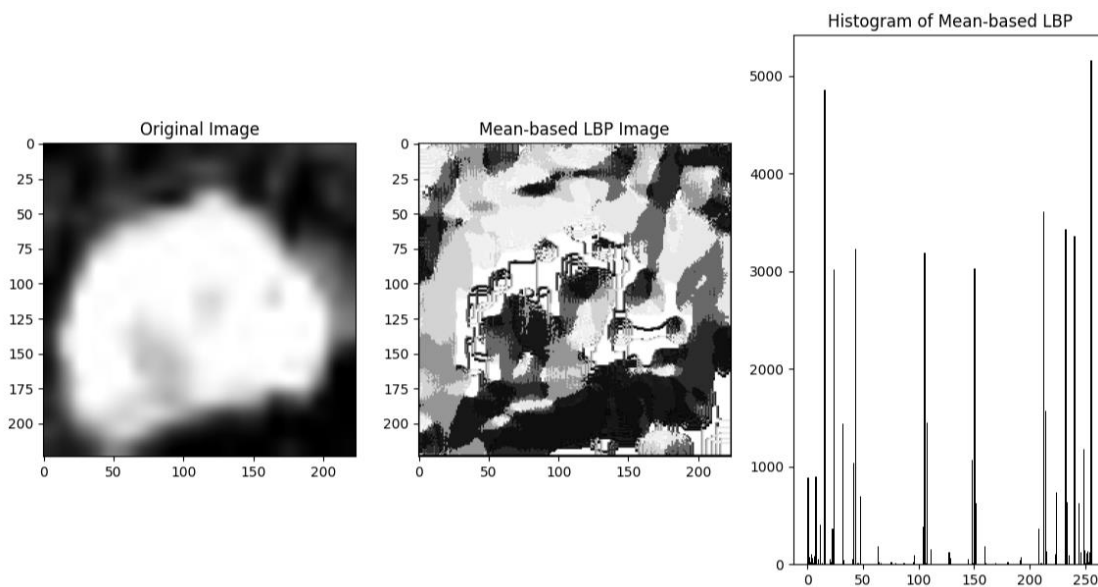
# Example usage
image_path = '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-2127/SUB7/p-049.jpg'

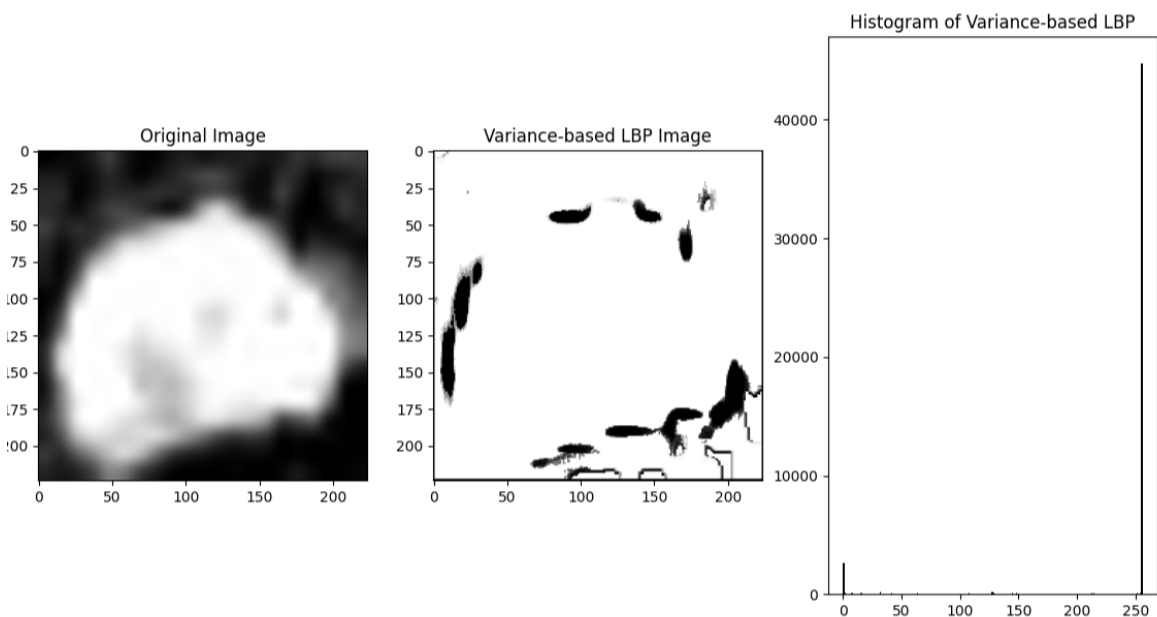

mean_based_lbp(image_path)

median_based_lbp(image_path)

variance_based_lbp(image_path)
```

Output:-





Now we will observe the mean median and variance based on 5 images

Code:-

```
import numpy as np

import matplotlib.pyplot as plt

from PIL import Image

def mean_based_lbp(img_array):

    rows, cols = img_array.shape

    lbp_image = np.zeros((rows, cols), dtype=np.uint8)

    for i in range(1, rows - 1):

        for j in range(1, cols - 1):

            neighborhood = img_array[i-1:i+2, j-1:j+2]

            center_pixel = img_array[i, j]

            mean_value = (np.sum(neighborhood) - center_pixel) / 8

            surrounding_pixels = np.delete(neighborhood.flatten(), 4)

            binary_pattern = (surrounding_pixels >=
mean_value).astype(int)

            lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))) .sum()

    return (lbp_image / lbp_image.max()) * 255
```

```

def median_based_lbp(img_array):

    rows, cols = img_array.shape

    lbp_image = np.zeros((rows, cols), dtype=np.uint8)

    for i in range(1, rows - 1):

        for j in range(1, cols - 1):

            neighborhood = img_array[i-1:i+2, j-1:j+2]

            center_pixel = img_array[i, j]

            median_value = np.median(np.delete(neighborhood.flatten(),
4))

            surrounding_pixels = np.delete(neighborhood.flatten(), 4)

            binary_pattern = (surrounding_pixels >=
median_value).astype(int)

            lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))).sum()

        return (lbp_image / lbp_image.max()) * 255

def variance_based_lbp(img_array):

    rows, cols = img_array.shape

    lbp_image = np.zeros((rows, cols), dtype=np.uint8)

    for i in range(1, rows - 1):

        for j in range(1, cols - 1):

            neighborhood = img_array[i-1:i+2, j-1:j+2]

            center_pixel = img_array[i, j]

```

```

        variance_value = np.var(np.delete(neighborhood.flatten(),
4))

        surrounding_pixels = np.delete(neighborhood.flatten(), 4)

        binary_pattern = (surrounding_pixels >=
variance_value).astype(int)

        lbp_image[i, j] = (binary_pattern * (2 **
np.arange(8))).sum()

    return (lbp_image / lbp_image.max()) * 255

# List of image paths
image_paths = [

    '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-2127/SUB7/p-
049.jpg',

    '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-1815/SUB3/p-
047.jpg',

    '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-1815/SUB3/p-
048.jpg',

    '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-1815/SUB3/p-
049.jpg',

    '/kaggle/input/bctumor/train/Malignant/BreaDM-Ma-1815/SUB3/p-
050.jpg'

]

# Processing and displaying all images
for idx, image_path in enumerate(image_paths):

    image = Image.open(image_path).convert('L').resize((224, 224))

    img_array = np.array(image)

```

```
mean_lbp = mean_based_lbp(img_array)

median_lbp = median_based_lbp(img_array)

variance_lbp = variance_based_lbp(img_array)


plt.figure(figsize=(15, 5))


plt.subplot(1, 4, 1)

plt.imshow(img_array, cmap='gray')

plt.title(f'Original Image {idx+1}')


plt.subplot(1, 4, 2)

plt.imshow(mean_lbp, cmap='gray')

plt.title('Mean-based LBP')


plt.subplot(1, 4, 3)

plt.imshow(median_lbp, cmap='gray')

plt.title('Median-based LBP')



plt.subplot(1, 4, 4)

plt.imshow(variance_lbp, cmap='gray')

plt.title('Variance-based LBP')


plt.show()
```


Output:-



GLCM Matrix Calculation:-

```
import numpy as np

def calculate_horizontal_glcmm(matrix, distance=1):

    max_gray_level = np.max(matrix) + 1

    glcmm = np.zeros((max_gray_level, max_gray_level), dtype=int)

    rows, cols = matrix.shape

    for i in range(rows):

        for j in range(cols - distance):

            current_pixel = matrix[i, j]

            neighbor_pixel = matrix[i, j + distance]

            glcmm[current_pixel, neighbor_pixel] += 1

    return glcmm

# Example input matrix
input_matrix = np.array([

    [0, 1, 2, 3],

    [1, 0, 3, 2],

    [2, 3, 0, 1],
```

```
[3, 2, 1, 0]

])

# Calculate the horizontal GLCM for the given input matrix

glcm_matrix = calculate_horizontal_glcm(input_matrix, distance=1)

print("Horizontal GLCM Matrix:\n", glcm_matrix)
```

output:-

```
Horizontal GLCM Matrix:
[[0 2 0 1]
 [2 0 1 0]
 [0 1 0 2]
 [1 0 2 0]]
```

Now we will start the model training part

Loading vgg16 model

```
import torch

import torchvision.models as models

vgg16 = models.vgg16(pretrained=True)

vgg16.eval()

print(vgg16)

/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing 'weight
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.cache/torch/hub/checkpoints/vgg16-397923af.pth
100%|#####| 528M/528M [00:10<00:00, 54.4MB/s]
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

Implementing a Custom Vgg16 function

Code:-

```
import torch

import torch.nn as nn

from torchvision import models

class CustomVGG16(nn.Module):

    def __init__(self, num_classes=2):

        # Initialize the parent class

        super(CustomVGG16, self).__init__()

        # Load the pre-trained VGG16 model

        vgg16 = models.vgg16(pretrained=True)
```

```

# Extract the features and avgpool layers

self.features = vgg16.features

self.avgpool = vgg16.avgpool


# Define a new classifier using nn.Sequential

self.classifier = nn.Sequential(

    nn.Linear(512 * 7 * 7, 4096), # First linear layer

    nn.ReLU(),                  # ReLU activation

    nn.Dropout(),               # Dropout layer

    nn.Linear(4096, 4096),      # Second linear layer

    nn.ReLU(),                  # ReLU activation

    nn.Dropout(),               # Dropout layer

    nn.Linear(4096, num_classes) # Final linear layer for
binary classification

)


def forward(self, x):

    # Pass the input through the features layer

    x = self.features(x)


    # Use the avgpool layer and reshape the output to a 2D tensor

    x = self.avgpool(x)


    x = torch.flatten(x, 1) # Flatten the tensor (batch_size,
num_features)


    # Pass the reshaped output to the custom classifier

```

```
x = self.classifier(x)
```

```
return x
```

```
# Example of creating an instance of CustomVGG16
```

```
model = CustomVGG16(num_classes=2)
```

Now, we will implement early stopping function

Code:-

```
import numpy as np

import torch

class EarlyStopping:

    def __init__(self, patience=7, verbose=False, delta=0,
path='checkpoint.pt', trace_func=print):

        """

        Early stops the training if validation loss doesn't improve
after a given patience.

        Parameters:

        - patience (int): How long to wait after the last time
validation loss improved.

        - verbose (bool): If True, prints a message for each validation
loss improvement.

        - delta (float): Minimum change in the monitored quantity to
qualify as an improvement.

        - path (str): Path for saving the model checkpoint.

        - trace_func (function): Function to print messages; can be set
to print/logging.

        """

        self.patience = patience

        self.verbose = verbose

        self.delta = delta

        self.path = path
```

```

self.trace_func = trace_func

self.counter = 0

self.best_score = None

self.early_stop = False

self.val_loss_min = np.Inf


def __call__(self, val_loss, model):
    """
    Call this function to check if the validation loss has
    improved.

    If not improved, increment the counter. If improved, reset
    counter and save model.

    Parameters:
    - val_loss (float): Current validation loss.
    - model (torch.nn.Module): Model to save if validation loss
    decreases.
    """
    score = -val_loss

    if self.best_score is None:
        self.best_score = score
        self.save_checkpoint(val_loss, model)
    elif score < self.best_score + self.delta:
        self.counter += 1

        self.trace_func(f'EarlyStopping counter: {self.counter} out
of {self.patience}')

        if self.counter >= self.patience:

```



```
        self.early_stop = True

    else:

        self.best_score = score

        self.save_checkpoint(val_loss, model)

        self.counter = 0

def save_checkpoint(self, val_loss, model):

    """Saves model when validation loss decreases."""

    if self.verbose:

        self.trace_func(f'Validation loss decreased  
({self.val_loss_min:.6f} --> {val_loss:.6f}). Saving model ...')

        torch.save(model.state_dict(), self.path)

        self.val_loss_min = val_loss
```

Now we will transform our images to make input in the model

```
_test_transforms = transforms.Compose([
    transforms.Resize(224),          # Resize the image to
    256x256
    transforms.CenterCrop(224),      # Center crop to 224x224
    transforms.ToTensor(),           # Convert to tensor
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
    0.225]) # Normalize with ImageNet mean and std
])
```

#Loading the Dataset

```
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

# Define transformations (resize and normalization only, no
augmentation)
transform = transforms.Compose([
    transforms.Resize((224, 224)),    # Resize the image to
    224x224
    transforms.ToTensor(),            # Convert the image to a
    PyTorch tensor
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
    0.225]) # Normalize based on ImageNet
])

# Replace BreastCancerDataset with ImageFolder for loading datasets
train_dir = '/kaggle/input/bctumor/train'
val_dir = '/kaggle/input/bctumor/val'
test_dir = '/kaggle/input/bctumor/test'

# Load datasets using ImageFolder
train_dataset = datasets.ImageFolder(root=train_dir,
transform=transform)
val_dataset = datasets.ImageFolder(root=val_dir, transform=transform)
test_dataset = datasets.ImageFolder(root=test_dir, transform=transform)

# Create DataLoaders for each dataset
```

```
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

# Now your datasets are ready for use
```

// training vgg16 model

```
import torch

import torch.nn.functional as F

from tqdm import tqdm

import torch.optim as optim

from tqdm import tqdm

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

model = model.to(device)
```

Train Function

```
import torch
from tqdm import tqdm
import torch.nn as nn
import torch.nn.functional as F

epoch = 0
total_epochs = 50
loader = train_loader # Ensure this is a DataLoader instance for
training data
criterion = nn.CrossEntropyLoss()
l2_decay = 0.01
lr = 0.01 # Learning rate
```

```

def train(epoch, model, total_epochs, loader, criterion, l2_decay):
    learning_rate = max(lr * (0.1 ** (epoch // 10)), 1e-5)
    optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate,
momentum=0.9, weight_decay=l2_decay)
    # Rest of the code...

    model.train()

    correct = 0
    for data, label in tqdm(loader, desc=f'Epoch
{epoch+1}/{total_epochs}', unit='batch'):
        #data = data.float().cuda()

        #label = label.long().cuda()
        data = data.float().to(device)
        label = label.long().to(device)

        output = model(data)
        optimizer.zero_grad()
        loss = F.nll_loss(F.log_softmax(output, dim=1), label)
        loss.backward()
        optimizer.step()

        pred = output.data.max(1)[1]
        correct += pred.eq(label.data.view_as(pred)).cpu().sum()

    print(f'train accuracy: {100. * correct / len(loader.dataset)}%')

```

#Validation Function

```

import torch
import numpy as np
from sklearn import metrics
from sklearn.metrics import roc_curve, auc as compute_auc
import torch.nn.functional as F

def validation(model, val_loader):
    model.eval() # Set model to evaluation mode
    test_loss = 0
    correct = 0
    all_predictions = [] # Store all predictions
    all_targets = [] # Store all targets
    possibilities = None # Store probabilities for AUC

```

```

for data, target in val_loader:
    if torch.cuda.is_available():
        data, target = data.cuda(), target.cuda()

    val_output = model(data)

    # Calculate test loss
    test_loss += F.nll_loss(F.log_softmax(val_output, dim=1),
target, reduction='sum').item()

    # Get predictions and accumulate them
    pred = val_output.data.max(1)[1]
    all_predictions.extend(pred.cpu().numpy()) # Collect all
predictions
    all_targets.extend(target.cpu().numpy()) # Collect all
target labels

    # Calculate probabilities for AUC
    possibility = F.softmax(val_output,
dim=1).cpu().detach().numpy()
    if possibilities is None:
        possibilities = possibility
    else:
        possibilities = np.concatenate((possibilities,
possibility), axis=0)

    # Calculate the number of correct predictions
    correct += pred.eq(target.data.view_as(pred)).cpu().sum()

    # Compute confusion matrix
    cm = metrics.confusion_matrix(all_targets, all_predictions)

    # One-hot encode the labels for AUC computation
    num_classes = val_output.shape[1]
    label_onehot =
np.eye(num_classes)[np.array(all_targets).astype(int)]

    # Compute ROC curve and AUC
    fpr, tpr, thresholds = roc_curve(label_onehot.ravel(),
possibilities.ravel())
    auc_value = compute_auc(fpr, tpr) # Use renamed function here

    # Average test loss per sample
    test_loss /= len(val_loader.dataset)

    # Calculate sensitivity and specificity
    specificity = 1 - fpr[1] if len(fpr) > 1 else 0

```

```
sensitivity = tpr[1] if len(tpr) > 1 else 0

print('Specificity: {:.4f}, Sensitivity: {:.4f}, AUC:
{:.4f}'.format(specificity, sensitivity, auc_value))
print('\nTest set: Average loss: {:.4f}, Accuracy:
{:.2f}%\n'.format(test_loss, 100. * correct / len(val_loader.dataset)))

return test_loss, 100. * correct / len(val_loader.dataset), cm,
auc_value
```

Now, Initiallizing the parameters and loading custom Vgg16

```

total_epochs = 50

lr = 0.01

momentum = 0.9

no_cuda = False

num_classes=2

log_interval = 10

l2_decay = 0.01

model = CustomVGG16(num_classes=num_classes)

model = model.to(device)

CustomVGG16(num_classes=num_classes)

```

```

/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.

```

```

    warnings.warn(
/opt/conda/lib/python3.10/site-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current
behavior is equivalent to passing `weights=VGG16_Weights.IMAGENET1K_V1`.
You can also use `weights=VGG16_Weights.DEFAULT` to get the most up-to-date
weights.

```

```

    warnings.warn(msg)
CustomVGG16(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
  )
)

```

```

        (13): ReLU(inplace=True)
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (18): ReLU(inplace=True)
        (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (20): ReLU(inplace=True)
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (22): ReLU(inplace=True)
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (25): ReLU(inplace=True)
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (27): ReLU(inplace=True)
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (29): ReLU(inplace=True)
        (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
    (classifier): Sequential(
      (0): Linear(in_features=25088, out_features=4096, bias=True)
      (1): ReLU()
      (2): Dropout(p=0.5, inplace=False)
      (3): Linear(in_features=4096, out_features=4096, bias=True)
      (4): ReLU()
      (5): Dropout(p=0.5, inplace=False)
      (6): Linear(in_features=4096, out_features=2, bias=True)
    )
  )
)

```

Now training the model

```
# please make sure you are using variables that match your own
environment.
```



```

# model training on Validation

model.to(device) # here device is cuda
best_accuracy = 0
early_stop = EarlyStopping(patience=20, verbose=True)
project_name = 'tumor_classification'
model_name = 'vgg16'

# we will be using epochs. epochs will be defined in another code
block.

for epoch in range(1, total_epochs + 1):

    #train(epoch, model)#train(epoch, total_epochs, train_loader,
criterion, l2_decay, lr)

    train(epoch, model, total_epochs, train_loader, criterion,
l2_decay)

    with torch.no_grad():

        test_loss, _, auc = validation(model, val_loader)

        #accuracy, test_loss, auc_value = test(model, test_loader)
        #print(f"Test Accuracy: {accuracy:.2f}%, Loss: {test_loss:.4f},
AUC: {auc_value:.4f}")

    # making sure that the model can run on multiple GPUs

    dict = model.module.state_dict() if isinstance(model,
nn.parallel.DistributedDataParallel) else model.state_dict()

    model_save_dir = os.path.join('model', project_name, model_name)

    if not os.path.exists(model_save_dir):

        os.makedirs(model_save_dir)

    early_stop(test_loss, model)

    if auc > best_accuracy:

        best_accuracy = auc

        #torch.save(os.path.join(model_save_dir,
f'{model_name}_{epoch}.pth'), _use_new_zipfile_serialization=False)
        torch.save(dict, os.path.join(model_save_dir,
f'{model_name}_{epoch}.pth'), _use_new_zipfile_serialization=False)

```

```
if early_stop.early_stop:

    print("Early stopping")

    break
```

Epoch 2/50: 100%|██████████| 639/639 [04:50<00:00, 2.20batch/s]
train accuracy: 84.44259643554688%
Specificity: 1.0000, Sensitivity: 0.0101, AUC: 0.9192

Test set: Average loss: 0.3461, Accuracy: 81.30%

Validation loss decreased (inf --> 0.346125). Saving model ...
Epoch 3/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 90.40814208984375%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8745

Test set: Average loss: 0.4374, Accuracy: 76.57%

EarlyStopping counter: 1 out of 20
Epoch 4/50: 100%|██████████| 639/639 [03:16<00:00, 3.25batch/s]
train accuracy: 90.92198944091797%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8659

Test set: Average loss: 0.4618, Accuracy: 76.02%

EarlyStopping counter: 2 out of 20
Epoch 5/50: 100%|██████████| 639/639 [03:20<00:00, 3.19batch/s]
train accuracy: 92.33140563964844%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8730

Test set: Average loss: 0.4737, Accuracy: 76.72%

EarlyStopping counter: 3 out of 20
Epoch 6/50: 100%|██████████| 639/639 [03:21<00:00, 3.18batch/s]
train accuracy: 92.71312713623047%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9151

Test set: Average loss: 0.3882, Accuracy: 82.00%

EarlyStopping counter: 4 out of 20
Epoch 7/50: 100%|██████████| 639/639 [03:16<00:00, 3.25batch/s]
train accuracy: 93.69677734375%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8963

Test set: Average loss: 0.4005, Accuracy: 79.14%

EarlyStopping counter: 5 out of 20
Epoch 8/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 94.07360076904297%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8835

Test set: Average loss: 0.4601, Accuracy: 78.48%

EarlyStopping counter: 6 out of 20
Epoch 9/50: 100%|██████████| 639/639 [03:17<00:00, 3.23batch/s]

train accuracy: 94.42595672607422%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8978

Test set: Average loss: 0.4009, Accuracy: 80.19%

EarlyStopping counter: 7 out of 20
Epoch 10/50: 100%|██████████| 639/639 [03:18<00:00, 3.23batch/s]
train accuracy: 94.48468017578125%
Specificity: 1.0000, Sensitivity: 0.1217, AUC: 0.9204

Test set: Average loss: 0.5308, Accuracy: 82.25%

EarlyStopping counter: 8 out of 20
Epoch 11/50: 100%|██████████| 639/639 [03:17<00:00, 3.23batch/s]
train accuracy: 98.82059478759766%
Specificity: 1.0000, Sensitivity: 0.0804, AUC: 0.9286

Test set: Average loss: 0.4507, Accuracy: 83.61%

EarlyStopping counter: 9 out of 20
Epoch 12/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 99.47636413574219%
Specificity: 1.0000, Sensitivity: 0.1111, AUC: 0.9336

Test set: Average loss: 0.4754, Accuracy: 84.26%

EarlyStopping counter: 10 out of 20
Epoch 13/50: 100%|██████████| 639/639 [03:16<00:00, 3.25batch/s]
train accuracy: 99.65743255615234%
Specificity: 1.0000, Sensitivity: 0.0870, AUC: 0.9251

Test set: Average loss: 0.4974, Accuracy: 83.11%

EarlyStopping counter: 11 out of 20
Epoch 14/50: 100%|██████████| 639/639 [03:16<00:00, 3.25batch/s]
train accuracy: 99.75531005859375%
Specificity: 1.0000, Sensitivity: 0.1212, AUC: 0.9164

Test set: Average loss: 0.5832, Accuracy: 82.35%

EarlyStopping counter: 12 out of 20
Epoch 15/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 99.76020050048828%
Specificity: 1.0000, Sensitivity: 0.1202, AUC: 0.9276

Test set: Average loss: 0.5303, Accuracy: 83.51%

EarlyStopping counter: 13 out of 20
Epoch 16/50: 100%|██████████| 639/639 [03:19<00:00, 3.20batch/s]
train accuracy: 99.75041961669922%
Specificity: 1.0000, Sensitivity: 0.1347, AUC: 0.9261

Test set: Average loss: 0.5683, Accuracy: 83.01%

EarlyStopping counter: 14 out of 20
Epoch 17/50: 100%|██████████| 639/639 [03:18<00:00, 3.22batch/s]
train accuracy: 99.72105407714844%
Specificity: 1.0000, Sensitivity: 0.0679, AUC: 0.8993

Test set: Average loss: 0.6361, Accuracy: 80.19%

EarlyStopping counter: 15 out of 20
Epoch 18/50: 100%|██████████| 639/639 [03:16<00:00, 3.26batch/s]
train accuracy: 99.7944564819336%
Specificity: 1.0000, Sensitivity: 0.1166, AUC: 0.9221

Test set: Average loss: 0.5539, Accuracy: 83.06%

EarlyStopping counter: 16 out of 20
Epoch 19/50: 100%|██████████| 639/639 [03:15<00:00, 3.27batch/s]
train accuracy: 99.77488708496094%
Specificity: 1.0000, Sensitivity: 0.1187, AUC: 0.9176

Test set: Average loss: 0.6069, Accuracy: 82.30%

EarlyStopping counter: 17 out of 20
Epoch 20/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 99.76020050048828%
Specificity: 1.0000, Sensitivity: 0.0744, AUC: 0.9182

Test set: Average loss: 0.5491, Accuracy: 83.16%

EarlyStopping counter: 18 out of 20
Epoch 21/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 99.9363784790039%
Specificity: 1.0000, Sensitivity: 0.1071, AUC: 0.9180

Test set: Average loss: 0.5881, Accuracy: 82.05%

EarlyStopping counter: 19 out of 20
Epoch 22/50: 100%|██████████| 639/639 [03:17<00:00, 3.24batch/s]
train accuracy: 99.97063446044922%
Specificity: 1.0000, Sensitivity: 0.1131, AUC: 0.9181

Test set: Average loss: 0.6001, Accuracy: 82.20%

EarlyStopping counter: 20 out of 20
Early stopping

#Test Function

```

from sklearn.metrics import roc_auc_score, roc_curve,
classification_report, confusion_matrix
import torch
import numpy as np
import torch.nn.functional as F

def test(model, test_loader):
    """
    Function to evaluate the model on the test dataset.

    Parameters:
    - model (torch.nn.Module): The trained model.
    - test_loader (DataLoader): DataLoader for the test dataset.

    Returns:
    - accuracy (float): Test accuracy in percentage.
    - test_loss (float): Average test loss.
    - auc_value (float): AUC score.
    """
    model.eval() # Set model to evaluation mode
    test_loss = 0
    correct = 0
    possibilities = None
    all_predictions = []
    all_targets = []

    # Iterate through the test DataLoader
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        test_output = model(data) # Forward pass

        # Compute loss
        test_loss += F.nll_loss(F.log_softmax(test_output, dim=1),
target, reduction='sum').item()

        # Get predictions
        pred = test_output.data.max(1)[1]
        all_predictions.extend(pred.cpu().numpy())
        all_targets.extend(target.cpu().numpy())

        # Get probabilities for AUC calculation
        possibility = F.softmax(test_output, dim=1).cpu().data.numpy()
        if possibilities is None:
            possibilities = possibility
        else:
            possibilities = np.concatenate((possibilities,
possibility), axis=0)

```

```

    # Count correct predictions
    correct += pred.eq(target.data.view_as(pred)).sum().item()

    # Convert all predictions and targets to numpy arrays
    all_predictions = np.array(all_predictions)
    all_targets = np.array(all_targets)

    # Classification metrics
    print(classification_report(all_targets, all_predictions,
target_names=['benign', 'malignant'], digits=4))

    # Confusion matrix
    cm = confusion_matrix(all_targets, all_predictions)
    print("Confusion Matrix:\n", cm)

    # AUC score
    num_classes = possibilities.shape[1]
    label_onehot = np.eye(num_classes)[all_targets]
    fpr, tpr, _ = roc_curve(label_onehot.ravel(),
possibilities.ravel())
    auc_value = roc_auc_score(label_onehot, possibilities,
average="macro")

    # Average loss and accuracy
    test_loss /= len(test_loader.dataset)
    accuracy = 100. * correct / len(test_loader.dataset)
    print('Specificity: {:.4f}, Sensitivity: {:.4f}, AUC:
{:.4f}'.format(1 - fpr[0], tpr[0], auc_value))
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{}
({:.2f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset), accuracy))

    return accuracy, test_loss, auc_value, fpr, tpr, cm

```

Output:-

test(model, test_loader)



	precision	recall	f1-score	support
benign	0.5010	0.3695	0.4253	1938
malignant	0.7746	0.8549	0.8128	4913
accuracy			0.7176	6851
macro avg	0.6378	0.6122	0.6190	6851
weighted avg	0.6972	0.7176	0.7032	6851

Confusion Matrix:

```
[[ 716 1222]
```

```
 [ 713 4200]]
```

Specificity: 1.0000, Sensitivity: 0.0000, AUC: 0.6716

Test set: Average loss: 1.7644, Accuracy: 4916/6851 (71.76%)

```
(71.75594803678295,  
 1.7643584101601721,  
 0.6716401243347351,  
 array([0.          , 0.00525471, 0.00919574, ..., 0.98963655, 0.98963655,  
        1.          ]),  
 array([0.          , 0.0507955 , 0.0719603 , ..., 0.99985404, 1.          ,  
        1.          ]),  
 array([[ 716, 1222],  
        [ 713, 4200]]))
```

test_accuracy, test_loss, test_auc, fpr, tpr, cm = test(model, test_loader)



	precision	recall	f1-score	support
benign	0.5010	0.3695	0.4253	1938
malignant	0.7746	0.8549	0.8128	4913
accuracy			0.7176	6851
macro avg	0.6378	0.6122	0.6190	6851
weighted avg	0.6972	0.7176	0.7032	6851

Confusion Matrix:

```
[[ 716 1222]
```

```
 [ 713 4200]]
```

Specificity: 1.0000, Sensitivity: 0.0000, AUC: 0.6716


Test set: Average loss: 1.7644, Accuracy: 4916/6851 (71.76%)

#plotting graph and confusion Matrix

```
import matplotlib.pyplot as plt

def plot_auc(fpr, tpr, auc_value):
    plt.figure(figsize=(5, 4))
    plt.plot(fpr, tpr, color='blue', lw=2, label=f'AUC = {auc_value:.4f}')
    plt.plot([0, 1], [0, 1], color='gray', linestyle='--', lw=1) # Diagonal line
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve')
    plt.legend(loc='lower right')
    plt.grid(alpha=0.3)
    plt.show()

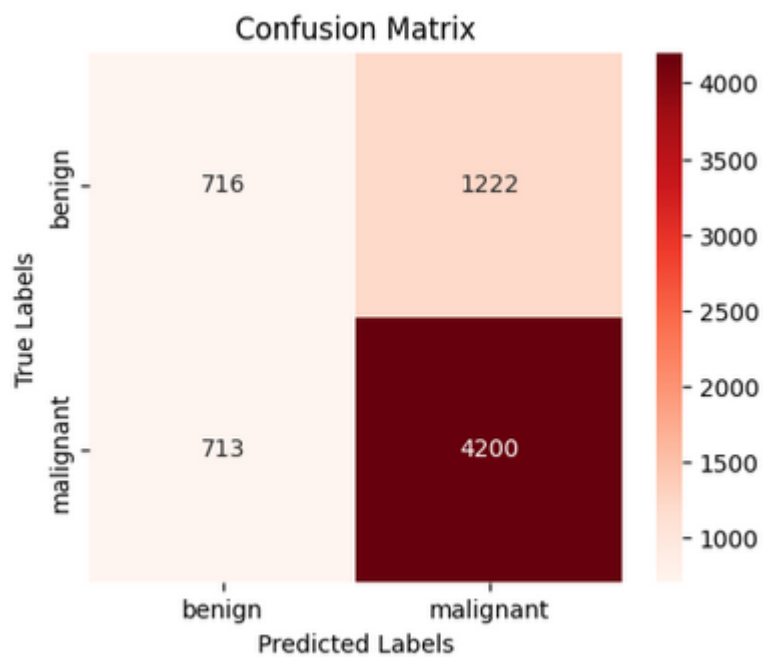
plot_auc(fpr, tpr, test_auc)
```



#Plotting Confusion Matrix

```
import seaborn as sns
def plot_confusion_matrix(cm, class_names):
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Reds',
xticklabels=class_names, yticklabels=class_names)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

plot_confusion_matrix(cm, class_names=['benign', 'malignant'])
```



#Implementing Resnet18

```
import torch
import torch.nn as nn
import torchvision.models as models

# ResNet18 class that inherits from nn.Module
class Resnet18(nn.Module):
    def __init__(self, num_classes=2):
        super(Resnet18, self).__init__()

        # Load the pretrained ResNet18 model
        model_resnet18 = models.resnet18(pretrained=True)

        # Extract layers from pretrained model
        self.conv1 = model_resnet18.conv1      # initial convolutional
layer
        self.bn1 = model_resnet18.bn1          # batch normalization
layer
        self.relu = model_resnet18.relu        # ReLU activation
function
        self.maxpool = model_resnet18.maxpool  # max pooling layer

        # ResNet blocks for feature extraction
        self.layer1 = model_resnet18.layer1
        self.layer2 = model_resnet18.layer2
        self.layer3 = model_resnet18.layer3
        self.layer4 = model_resnet18.layer4    # deeper layers for
increasing depth of the network

        # Average pooling layer
        self.avgpool = model_resnet18.avgpool

        # Replace the fully connected layer for custom number of
classes
        self._features = model_resnet18.fc.in_features
        self.fc = nn.Linear(self._features, num_classes)

    def forward(self, x):
        x = self.conv1(x)                      # apply convolutional
layer
        x = self.bn1(x)                        # apply batch
normalization
        x = self.relu(x)                      # apply ReLU activation
        x = self.maxpool(x)                   # apply max pooling
        x = self.layer1(x)                    # pass through ResNet
layer 1
```

```

        x = self.layer2(x)                # pass through ResNet
layer 2
        x = self.layer3(x)                # pass through ResNet
layer 3
        x = self.layer4(x)                # pass through ResNet
layer 4
        x = self.avgpool(x)               # apply average pooling
        x = x.view(x.size(0), -1)         # flatten for the fully
connected layer
        x = self.fc(x)                    # apply fully connected
layer for output
        return x

# Print the model architecture
if __name__ == "__main__":
    model1 = Resnet18(num_classes=2)
    #print(model1)

```

Now Training Resnet18

```
import torch.nn as nn
total_epochs = 50
lr = 0.01
momentum = 0.9
no_cuda = False
num_classes=2
log_interval = 10
l2_decay = 0.01
model = Resnet18(num_classes=num_classes)
model = model.to(device)
criterion = nn.CrossEntropyLoss()
```

```
from torchvision.models import resnet18

# Initialize ResNet18 model
model = resnet18(pretrained=True) # Load a pre-trained ResNet18
num_fts = model.fc.in_features    # Get the number of input features
to the fully connected layer
model.fc = nn.Linear(num_fts, num_classes) # Replace the fully
connected layer for your specific task
model = model.to(device)

# Loss function
criterion = nn.CrossEntropyLoss()

# Early stopping setup (implement EarlyStopping in your environment)
early_stop = EarlyStopping(patience=20, verbose=True)

# Project and model names
project_name = 'tumor_classification'
model_name = 'ResNet18'

best_accuracy = 0
model_save_dir = os.path.join('model', project_name, model_name)
os.makedirs(model_save_dir, exist_ok=True)

# Training loop
for epoch in range(1, total_epochs + 1):
    train(epoch, model, total_epochs, train_loader, criterion,
    l2_decay)
    with torch.no_grad():
        test_loss, accuracy, cm, auc = validation(model, val_loader)
```

```

# Save the model if it achieves the best AUC
model_dict = model.state_dict()
if auc > best_accuracy:
    best_accuracy = auc
    torch.save(model_dict, os.path.join(model_save_dir,
f'{model_name}_{epoch}.pth'))
    early_stop(test_loss, model)
if early_stop.early_stop:
    print("Early stopping")
    break

```

Output :-

Epoch 2/50: 100%|██████████| 639/639 [01:18<00:00, 8.10batch/s]
train accuracy: 87.0558853149414%
Specificity: 1.0000, Sensitivity: 0.0010, AUC: 0.8515

Test set: Average loss: 0.6394, Accuracy: 73.25%

Validation loss decreased (inf --> 0.639401). Saving model ...
Epoch 3/50: 100%|██████████| 639/639 [01:19<00:00, 8.05batch/s]
train accuracy: 94.05403137207031%
Specificity: 1.0000, Sensitivity: 0.0111, AUC: 0.9451

Test set: Average loss: 0.3697, Accuracy: 87.38%

Validation loss decreased (0.639401 --> 0.369708). Saving model ...
Epoch 4/50: 100%|██████████| 639/639 [01:20<00:00, 7.98batch/s]
train accuracy: 95.8255844116211%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9266

Test set: Average loss: 0.4239, Accuracy: 84.62%

EarlyStopping counter: 1 out of 20
Epoch 5/50: 100%|██████████| 639/639 [01:17<00:00, 8.29batch/s]
train accuracy: 96.45198822021484%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8545

Test set: Average loss: 0.5771, Accuracy: 77.68%

EarlyStopping counter: 2 out of 20
Epoch 6/50: 100%|██████████| 639/639 [01:14<00:00, 8.52batch/s]
train accuracy: 97.27904510498047%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8505

Test set: Average loss: 0.5468, Accuracy: 76.87%

EarlyStopping counter: 3 out of 20
Epoch 7/50: 100%|██████████| 639/639 [01:18<00:00, 8.17batch/s]
train accuracy: 97.2398910522461%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9143

Test set: Average loss: 0.4178, Accuracy: 81.75%

EarlyStopping counter: 4 out of 20
Epoch 8/50: 100%|██████████| 639/639 [01:19<00:00, 8.05batch/s]

train accuracy: 97.53841400146484%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8669

Test set: Average loss: 0.5115, Accuracy: 78.33%

EarlyStopping counter: 5 out of 20
Epoch 9/50: 100%|██████████| 639/639 [01:20<00:00, 7.95batch/s]
train accuracy: 97.72438049316406%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9117

Test set: Average loss: 0.4347, Accuracy: 83.86%

EarlyStopping counter: 6 out of 20
Epoch 10/50: 100%|██████████| 639/639 [01:20<00:00, 7.93batch/s]
train accuracy: 97.54820251464844%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9003

Test set: Average loss: 0.5119, Accuracy: 82.35%

EarlyStopping counter: 7 out of 20
Epoch 11/50: 100%|██████████| 639/639 [01:21<00:00, 7.85batch/s]
train accuracy: 99.73573303222656%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9021

Test set: Average loss: 0.4736, Accuracy: 82.96%

EarlyStopping counter: 8 out of 20
Epoch 12/50: 100%|██████████| 639/639 [01:17<00:00, 8.28batch/s]
train accuracy: 99.91680908203125%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9088

Test set: Average loss: 0.4493, Accuracy: 84.11%

EarlyStopping counter: 9 out of 20
Epoch 13/50: 100%|██████████| 639/639 [01:16<00:00, 8.33batch/s]
train accuracy: 99.96084594726562%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9180

Test set: Average loss: 0.4122, Accuracy: 85.17%

EarlyStopping counter: 10 out of 20
Epoch 14/50: 100%|██████████| 639/639 [01:16<00:00, 8.32batch/s]
train accuracy: 99.98042297363281%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.8931

Test set: Average loss: 0.4911, Accuracy: 83.06%

EarlyStopping counter: 11 out of 20
Epoch 15/50: 100%|██████████| 639/639 [01:13<00:00, 8.64batch/s]
train accuracy: 99.99510955810547%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.8958

Test set: Average loss: 0.4745, Accuracy: 82.55%

EarlyStopping counter: 12 out of 20
Epoch 16/50: 100%|██████████| 639/639 [01:14<00:00, 8.60batch/s]
train accuracy: 99.9902114868164%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9164

Test set: Average loss: 0.4133, Accuracy: 85.52%

EarlyStopping counter: 13 out of 20
Epoch 17/50: 100%|██████████| 639/639 [01:14<00:00, 8.60batch/s]
train accuracy: 100.0%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9111

Test set: Average loss: 0.4263, Accuracy: 84.26%

EarlyStopping counter: 14 out of 20
Epoch 18/50: 100%|██████████| 639/639 [01:14<00:00, 8.57batch/s]
train accuracy: 99.99510955810547%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9035

Test set: Average loss: 0.4541, Accuracy: 84.06%

EarlyStopping counter: 15 out of 20
Epoch 19/50: 100%|██████████| 639/639 [01:17<00:00, 8.29batch/s]
train accuracy: 99.9902114868164%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9151

Test set: Average loss: 0.4163, Accuracy: 85.37%

EarlyStopping counter: 16 out of 20
Epoch 20/50: 100%|██████████| 639/639 [01:17<00:00, 8.21batch/s]
train accuracy: 99.9902114868164%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9044

Test set: Average loss: 0.4416, Accuracy: 84.16%

EarlyStopping counter: 17 out of 20
Epoch 21/50: 100%|██████████| 639/639 [01:18<00:00, 8.14batch/s]
train accuracy: 100.0%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9131

Test set: Average loss: 0.4198, Accuracy: 84.97%

EarlyStopping counter: 18 out of 20
Epoch 22/50: 100%|██████████| 639/639 [01:20<00:00, 7.96batch/s]
train accuracy: 100.0%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9098

Test set: Average loss: 0.4310, Accuracy: 84.77%

EarlyStopping counter: 19 out of 20
Epoch 23/50: 100%|██████████| 639/639 [01:18<00:00, 8.16batch/s]
train accuracy: 99.99510955810547%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.9183

Test set: Average loss: 0.4066, Accuracy: 85.72%

EarlyStopping counter: 20 out of 20
Early stopping

Resnet18 Testing

```
from sklearn.metrics import roc_auc_score, roc_curve,
classification_report, confusion_matrix
import torch
import numpy as np
import torch.nn.functional as F

def test(model, test_loader):
    """
    Function to evaluate the model on the test dataset.

    Parameters:
    - model (torch.nn.Module): The trained model.
    - test_loader (DataLoader): DataLoader for the test dataset.

    Returns:
    - accuracy (float): Test accuracy in percentage.
    - test_loss (float): Average test loss.
    - auc_value (float): AUC score.
    """
    model.eval() # Set model to evaluation mode
    test_loss = 0
    correct = 0
    possibilities = None
    all_predictions = []
    all_targets = []

    # Iterate through the test DataLoader
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        test_output = model(data) # Forward pass

        # Compute loss
        test_loss += F.nll_loss(F.log_softmax(test_output, dim=1),
target, reduction='sum').item()

        # Get predictions
        pred = test_output.data.max(1)[1]
        all_predictions.extend(pred.cpu().numpy())
        all_targets.extend(target.cpu().numpy())

        # Get probabilities for AUC calculation
        possibility = F.softmax(test_output, dim=1).cpu().data.numpy()
        if possibilities is None:
            possibilities = possibility
        else:
```



```

        possibilities = np.concatenate((possibilities,
possibility), axis=0)

    # Count correct predictions
    correct += pred.eq(target.data.view_as(pred)).sum().item()

    # Convert all predictions and targets to numpy arrays
    all_predictions = np.array(all_predictions)
    all_targets = np.array(all_targets)

    # Classification metrics
    print(classification_report(all_targets, all_predictions,
target_names=['benign', 'malignant'], digits=4))

    # Confusion matrix
    cm = confusion_matrix(all_targets, all_predictions)
    print("Confusion Matrix:\n", cm)

    # AUC score
    num_classes = possibilities.shape[1]
    label_onehot = np.eye(num_classes)[all_targets]
    fpr, tpr, _ = roc_curve(label_onehot.ravel(),
possibilities.ravel())
    auc_value = roc_auc_score(label_onehot, possibilities,
average="macro")

    # Average loss and accuracy
    test_loss /= len(test_loader.dataset)
    accuracy = 100. * correct / len(test_loader.dataset)
    print('Specificity: {:.4f}, Sensitivity: {:.4f}, AUC:
{:.4f}'.format(1 - fpr[0], tpr[0], auc_value))
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{ }
({:.2f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset), accuracy))

    return accuracy, test_loss, auc_value, fpr, tpr

```

Output:-

```
▶ test_accuracy, test_loss, test_auc, fpr, tpr = test(model, test_loader)
```



	precision	recall	f1-score	support
benign	0.6200	0.2147	0.3189	1938
malignant	0.7537	0.9481	0.8398	4913
accuracy			0.7406	6851
macro avg	0.6868	0.5814	0.5794	6851
weighted avg	0.7159	0.7406	0.6925	6851

Confusion Matrix:

```
[[ 416 1522]
```

```
 [ 255 4658]]
```

Specificity: 1.0000, Sensitivity: 0.0000, AUC: 0.7190

Test set: Average loss: 0.8269, Accuracy: 5074/6851 (74.06%)


NOW Plotting graph for the same model :-

```

import matplotlib.pyplot as plt

def plot_auc(fpr, tpr, auc_value):
    plt.figure(figsize=(5, 4))
    plt.plot(fpr, tpr, color='blue', lw=2, label=f'AUC = {auc_value:.4f}')
    plt.plot([0, 1], [0, 1], color='gray', linestyle='--', lw=1) # Diagonal line
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve')
    plt.legend(loc='lower right')
    plt.grid(alpha=0.3)
    plt.show()
plot_auc(fpr, tpr, test_auc)

```




Now we will plot the confusion matrix

```
import seaborn as sns

def plot_confusion_matrix(cm, class_names):
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names, yticklabels=class_names)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()

plot_confusion_matrix(cm, class_names=['benign', 'malignant'])
```



NOW we will implement resnet 50 as same as resnet18

```
import torch
import torch.nn as nn
import torchvision.models as models

class resnet50(nn.Module):
    def __init__(self, num_classes=2):
        super(resnet50, self).__init__()

        # Load the pretrained ResNet50 model
        model_resnet50 = models.resnet50(pretrained=True)

        # Extract layers from the pretrained model
        self.conv1 = model_resnet50.conv1          # initial
convolutional layer
        self.bn1 = model_resnet50.bn1              # batch normalization
layer
        self.relu = model_resnet50.relu            # ReLU activation
function
        self.maxpool = model_resnet50.maxpool      # max pooling layer

        # ResNet blocks for feature extraction
        self.layer1 = model_resnet50.layer1
        self.layer2 = model_resnet50.layer2
        self.layer3 = model_resnet50.layer3
        self.layer4 = model_resnet50.layer4        # deeper layers for
increasing depth of the network

        # Average pooling layer
        self.avgpool = model_resnet50.avgpool

        # Replace the fully connected layer for custom number of
classes
        self._features = model_resnet50.fc.in_features
        self.fc = nn.Linear(self._features, num_classes)

    def forward(self, x):
        x = self.conv1(x)                          # apply convolutional
layer
        x = self.bn1(x)                            # apply batch
normalization
        x = self.relu(x)                          # apply ReLU
activation
        x = self.maxpool(x)                       # apply max pooling
```

```

        x = self.layer1(x)                # pass through ResNet
layer 1
        x = self.layer2(x)                # pass through ResNet
layer 2
        x = self.layer3(x)                # pass through ResNet
layer 3
        x = self.layer4(x)                # pass through ResNet
layer 4
        x = self.avgpool(x)               # apply average
pooling
        x = x.view(x.size(0), -1)         # flatten for the
fully connected layer
        x = self.fc(x)                    # apply fully
connected layer for output
        return x

# Print the model architecture
if __name__ == "__main__":
    model2 = resnet50(num_classes=2)
    #print(model2)

```

/// Initializing the parameters and training the model

```

import torch.nn as nn
total_epochs = 50
lr = 0.01
momentum = 0.9
no_cuda = False
num_classes=2
log_interval = 10
l2_decay = 0.01
model = resnet50(num_classes=num_classes)
model = model.to(device)
criterion = nn.CrossEntropyLoss()

```

```

from torchvision.models import resnet50

# Initialize ResNet50 model
model = resnet50(pretrained=True)  # Load a pre-trained ResNet50

```

```

num_ftrs = model.fc.in_features    # Get the number of input features
to the fully connected layer
model.fc = nn.Linear(num_ftrs, num_classes) # Replace the fully
connected layer for your specific task
model = model.to(device)

# Loss function
criterion = nn.CrossEntropyLoss()

# Early stopping setup (implement EarlyStopping in your environment)
early_stop = EarlyStopping(patience=20, verbose=True)

# Project and model names
project_name = 'tumor_classification'
model_name = 'ResNet50'

best_accuracy = 0
model_save_dir = os.path.join('model', project_name, model_name)
os.makedirs(model_save_dir, exist_ok=True)

# Training loop
for epoch in range(1, total_epochs + 1):
    train(epoch, model, total_epochs, train_loader, criterion,
    lr_decay)
    with torch.no_grad():
        test_loss, accuracy, cm, auc = validation(model, val_loader)
        # Save the model if it achieves the best AUC
        model_dict = model.state_dict()
        if auc > best_accuracy:
            best_accuracy = auc
            torch.save(model_dict, os.path.join(model_save_dir,
f'{model_name}_{epoch}.pth'))
            early_stop(test_loss, model)
            if early_stop.early_stop:
                print("Early stopping")
                break
        # Save the model at the final epoch if early stopping is triggered
        final_model_save_path = os.path.join(model_save_dir,
f'{model_name}_final_epoch.pth')
        torch.save(model.state_dict(), final_model_save_path,
        _use_new_zipfile_serialization=False)
        print(f"Final model saved at:
{os.path.abspath(final_model_save_path)}")

```

Output :-

Epoch 2/50: 100%|██████████| 639/639 [02:36<00:00, 4.08batch/s]
train accuracy: 86.9139633178711%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8513

Test set: Average loss: 0.5529, Accuracy: 75.11%

Validation loss decreased (inf --> 0.552876). Saving model ...
Epoch 3/50: 100%|██████████| 639/639 [02:35<00:00, 4.11batch/s]
train accuracy: 91.6952133178711%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8775

Test set: Average loss: 0.4440, Accuracy: 78.73%

Validation loss decreased (0.552876 --> 0.443954). Saving model ...
Epoch 4/50: 100%|██████████| 639/639 [02:31<00:00, 4.22batch/s]
train accuracy: 91.53372192382812%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8421

Test set: Average loss: 0.5380, Accuracy: 76.07%

EarlyStopping counter: 1 out of 20
Epoch 5/50: 100%|██████████| 639/639 [02:28<00:00, 4.29batch/s]
train accuracy: 92.8354721069336%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9130

Test set: Average loss: 0.3709, Accuracy: 81.50%

Validation loss decreased (0.443954 --> 0.370907). Saving model ...
Epoch 6/50: 100%|██████████| 639/639 [02:28<00:00, 4.29batch/s]
train accuracy: 93.51570892333984%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.6146

Test set: Average loss: 1.0713, Accuracy: 59.63%

EarlyStopping counter: 1 out of 20
Epoch 7/50: 100%|██████████| 639/639 [02:27<00:00, 4.34batch/s]
train accuracy: 94.16168975830078%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8875

Test set: Average loss: 0.5882, Accuracy: 79.39%

EarlyStopping counter: 2 out of 20
Epoch 8/50: 100%|██████████| 639/639 [02:27<00:00, 4.33batch/s]
train accuracy: 93.98062133789062%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8523

Test set: Average loss: 0.5926, Accuracy: 76.57%

EarlyStopping counter: 3 out of 20
Epoch 9/50: 100%|██████████| 639/639 [02:26<00:00, 4.35batch/s]
train accuracy: 94.79788208007812%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.9166

Test set: Average loss: 0.4207, Accuracy: 83.86%

EarlyStopping counter: 4 out of 20

Epoch 10/50: 100%|██████████| 639/639 [02:26<00:00, 4.36batch/s]
train accuracy: 94.91533660888672%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.7946

Test set: Average loss: 0.7814, Accuracy: 73.20%

EarlyStopping counter: 5 out of 20
Epoch 11/50: 100%|██████████| 639/639 [02:26<00:00, 4.37batch/s]
train accuracy: 98.99187469482422%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8335

Test set: Average loss: 0.6837, Accuracy: 74.26%

EarlyStopping counter: 6 out of 20
Epoch 12/50: 100%|██████████| 639/639 [02:27<00:00, 4.35batch/s]
train accuracy: 99.68679809570312%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8461

Test set: Average loss: 0.6570, Accuracy: 76.12%

EarlyStopping counter: 7 out of 20
Epoch 13/50: 100%|██████████| 639/639 [02:26<00:00, 4.35batch/s]
train accuracy: 99.82872009277344%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8291

Test set: Average loss: 0.7117, Accuracy: 74.51%

EarlyStopping counter: 8 out of 20
Epoch 14/50: 100%|██████████| 639/639 [02:29<00:00, 4.28batch/s]
train accuracy: 99.85318756103516%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8142

Test set: Average loss: 0.7894, Accuracy: 72.10%

EarlyStopping counter: 9 out of 20
Epoch 15/50: 100%|██████████| 639/639 [02:32<00:00, 4.19batch/s]
train accuracy: 99.89723205566406%
Specificity: 0.9995, Sensitivity: 0.0000, AUC: 0.8332

Test set: Average loss: 0.7160, Accuracy: 75.31%

EarlyStopping counter: 10 out of 20
Epoch 16/50: 100%|██████████| 639/639 [02:33<00:00, 4.16batch/s]
train accuracy: 99.90702056884766%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8394

Test set: Average loss: 0.6747, Accuracy: 75.92%

EarlyStopping counter: 11 out of 20
Epoch 17/50: 100%|██████████| 639/639 [02:28<00:00, 4.30batch/s]
train accuracy: 99.9363784790039%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8091

Test set: Average loss: 0.8079, Accuracy: 71.24%

EarlyStopping counter: 12 out of 20
Epoch 18/50: 100%|██████████| 639/639 [02:27<00:00, 4.33batch/s]
train accuracy: 99.91191101074219%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8396

Test set: Average loss: 0.6782, Accuracy: 76.52%

EarlyStopping counter: 13 out of 20
Epoch 19/50: 100%|██████████| 639/639 [02:27<00:00, 4.33batch/s]
train accuracy: 99.91680908203125%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.7977

Test set: Average loss: 0.8371, Accuracy: 70.04%

EarlyStopping counter: 14 out of 20
Epoch 20/50: 100%|██████████| 639/639 [02:27<00:00, 4.32batch/s]
train accuracy: 99.81403350830078%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8087

Test set: Average loss: 0.7207, Accuracy: 71.54%

EarlyStopping counter: 15 out of 20
Epoch 21/50: 100%|██████████| 639/639 [02:27<00:00, 4.34batch/s]
train accuracy: 99.9461669921875%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8558

Test set: Average loss: 0.6294, Accuracy: 75.52%

EarlyStopping counter: 16 out of 20
Epoch 22/50: 100%|██████████| 639/639 [02:28<00:00, 4.32batch/s]
train accuracy: 99.97553253173828%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8248

Test set: Average loss: 0.7349, Accuracy: 72.95%

EarlyStopping counter: 17 out of 20
Epoch 23/50: 100%|██████████| 639/639 [02:30<00:00, 4.26batch/s]
train accuracy: 99.97553253173828%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8447

Test set: Average loss: 0.6761, Accuracy: 75.01%

EarlyStopping counter: 18 out of 20
Epoch 24/50: 100%|██████████| 639/639 [02:27<00:00, 4.33batch/s]
train accuracy: 99.9559555053711%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8383

Test set: Average loss: 0.6953, Accuracy: 75.31%

EarlyStopping counter: 19 out of 20
Epoch 25/50: 100%|██████████| 639/639 [02:29<00:00, 4.28batch/s]
train accuracy: 99.96574401855469%
Specificity: 1.0000, Sensitivity: 0.0005, AUC: 0.8217

Test set: Average loss: 0.7503, Accuracy: 72.90%

EarlyStopping counter: 20 out of 20
Early stopping
Final model saved at:
/kaggle/working/model/tumor_classification/ResNet50/ResNet50_final_epoch.pt
h

// Implementing Testing Function

```
from sklearn.metrics import roc_auc_score, roc_curve,
classification_report, confusion_matrix
import torch
import numpy as np
import torch.nn.functional as F

def test(model, test_loader):
    """
    Function to evaluate the model on the test dataset.

    Parameters:
    - model (torch.nn.Module): The trained model.
    - test_loader (DataLoader): DataLoader for the test dataset.

    Returns:
    - accuracy (float): Test accuracy in percentage.
    - test_loss (float): Average test loss.
    - auc_value (float): AUC score.
    """
    model.eval() # Set model to evaluation mode
    test_loss = 0
    correct = 0
    possibilities = None
    all_predictions = []
    all_targets = []

    # Iterate through the test DataLoader
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        test_output = model(data) # Forward pass

        # Compute loss
        test_loss += F.nll_loss(F.log_softmax(test_output, dim=1),
target, reduction='sum').item()

        # Get predictions
        pred = test_output.data.max(1)[1]
        all_predictions.extend(pred.cpu().numpy())
        all_targets.extend(target.cpu().numpy())

        # Get probabilities for AUC calculation
        possibility = F.softmax(test_output, dim=1).cpu().data.numpy()
        if possibilities is None:
            possibilities = possibility
        else:
```

```

        possibilities = np.concatenate((possibilities,
possibility), axis=0)

    # Count correct predictions
    correct += pred.eq(target.data.view_as(pred)).sum().item()

    # Convert all predictions and targets to numpy arrays
    all_predictions = np.array(all_predictions)
    all_targets = np.array(all_targets)

    # Classification metrics
    print(classification_report(all_targets, all_predictions,
target_names=['benign', 'malignant'], digits=4))

    # Confusion matrix
    cm = confusion_matrix(all_targets, all_predictions)
    print("Confusion Matrix:\n", cm)

    # AUC score
    num_classes = possibilities.shape[1]
    label_onehot = np.eye(num_classes)[all_targets]
    fpr, tpr, _ = roc_curve(label_onehot.ravel(),
possibilities.ravel())
    auc_value = roc_auc_score(label_onehot, possibilities,
average="macro")

    # Average loss and accuracy
    test_loss /= len(test_loader.dataset)
    accuracy = 100. * correct / len(test_loader.dataset)
    print('Specificity: {:.4f}, Sensitivity: {:.4f}, AUC:
{:.4f}'.format(1 - fpr[0], tpr[0], auc_value))
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{ }
({:.2f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset), accuracy))

    return test_loss, accuracy, cm, auc_value

```

Output:-

resnet50 testing

```
[ ] test_accuracy, test_loss, test_auc, fpr = test(model, test_loader)
```

```
↔
      precision    recall  f1-score   support

   benign      0.5520      0.3617      0.4370      1938
  malignant      0.7784      0.8842      0.8279      4913

   accuracy                   0.7364      6851
  macro avg      0.6652      0.6229      0.6325      6851
 weighted avg      0.7143      0.7364      0.7173      6851
```

Confusion Matrix:

```
[[ 701 1237]
 [ 569 4344]]
```


Specificity: 1.0000, Sensitivity: 0.0000, AUC: 0.7018

Test set: Average loss: 0.8197, Accuracy: 5045/6851 (73.64%)

Now plotting graph and confusion matrix


```
import matplotlib.pyplot as plt

def plot_auc(fpr, tpr, auc_value):
    plt.figure(figsize=(5, 4))
    plt.plot(fpr, tpr, color='blue', lw=2, label=f'AUC =
{auc_value:.4f}')
    plt.plot([0, 1], [0, 1], color='gray', linestyle='--', lw=1) #
Diagonal line
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) Curve')
    plt.legend(loc='lower right')
    plt.grid(alpha=0.3)
    plt.show()
plot_auc(fpr, tpr, test_auc)
```



```
import seaborn as sns

def plot_confusion_matrix(cm, class_names):
    plt.figure(figsize=(5, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names)
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.title('Confusion Matrix')
    plt.show()
plot_confusion_matrix(cm, class_names=['benign', 'malignant'])
```



NOW with all the code we will implement Gradio for UI for our Models

```
[ ] pip install gradio
```

```
Collecting gradio
  Downloading gradio-5.6.0-py3-none-any.whl.metadata (16 kB)
Requirement already satisfied: aiofiles<24.0,>=22.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (22.1.0)
Requirement already satisfied: anyio<5.0,>=3.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (4.4.0)
Collecting fastapi<1.0,>=0.115.2 (from gradio)
  Downloading fastapi-0.115.5-py3-none-any.whl.metadata (27 kB)
Collecting ffmpeg (from gradio)
  Downloading ffmpeg-0.4.0-py3-none-any.whl.metadata (2.9 kB)
Collecting gradio-client==1.4.3 (from gradio)
  Downloading gradio_client-1.4.3-py3-none-any.whl.metadata (7.1 kB)
Requirement already satisfied: httpx>=0.24.1 in /opt/conda/lib/python3.10/site-packages (from gradio) (0.27.0)
Requirement already satisfied: huggingface-hub>=0.25.1 in /opt/conda/lib/python3.10/site-packages (from gradio) (0.25.1)
Requirement already satisfied: Jinja2<4.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (3.1.4)
Requirement already satisfied: MarkupSafe==2.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (2.1.5)
Requirement already satisfied: numpy<3.0,>=1.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (1.26.4)
Requirement already satisfied: orjson==3.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (3.10.4)
Requirement already satisfied: packaging in /opt/conda/lib/python3.10/site-packages (from gradio) (21.3)
Requirement already satisfied: pandas<3.0,>=1.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (2.2.2)
Requirement already satisfied: pillow<12.0,>=8.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (10.3.0)
Requirement already satisfied: pydantic>=2.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (2.9.2)
Requirement already satisfied: pydub in /opt/conda/lib/python3.10/site-packages (from gradio) (0.25.1)
Collecting python-multipart==0.0.12 (from gradio)
  Downloading python_multipart-0.0.12-py3-none-any.whl.metadata (1.9 kB)
Requirement already satisfied: pyyaml<7.0,>=5.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (6.0.2)
Collecting ruff>=0.2.2 (from gradio)
  Downloading ruff-0.8.0-py3-none-manylinux_2_17_x86_64_musllinux_2_17_x86_64.whl.metadata (25 kB)
Collecting safehttpx<1.0,>=0.1.1 (from gradio)
  Downloading safehttpx-0.1.1-py3-none-any.whl.metadata (4.1 kB)
Collecting semantic-version==2.0 (from gradio)
  Downloading semantic_version-2.10.0-py2.py3-none-any.whl.metadata (9.7 kB)
Collecting starlette<1.0,>=0.40.0 (from gradio)
  Downloading starlette-0.41.3-py3-none-any.whl.metadata (6.0 kB)
Collecting tomlkit==0.12.0 (from gradio)
  Downloading tomlkit-0.12.0-py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: typer<1.0,>=0.12 in /opt/conda/lib/python3.10/site-packages (from gradio) (0.12.3)
Requirement already satisfied: typing-extensions==4.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (4.12.2)
Requirement already satisfied: uvicorn>=0.14.0 in /opt/conda/lib/python3.10/site-packages (from gradio) (0.30.1)
Requirement already satisfied: fsspec in /opt/conda/lib/python3.10/site-packages (from gradio-client==1.4.3->gradio) (2024.6.1)
Requirement already satisfied: websockets<13.0,>=10.0 in /opt/conda/lib/python3.10/site-packages (from gradio-client==1.4.3->gradio) (12.0)
Requirement already satisfied: idna>=2.8 in /opt/conda/lib/python3.10/site-packages (from anyio<5.0,>=3.0->gradio) (3.7)
Requirement already satisfied: sniffio==1.1 in /opt/conda/lib/python3.10/site-packages (from anyio<5.0,>=3.0->gradio) (1.3.1)
Requirement already satisfied: exceptiongroup>=1.0.2 in /opt/conda/lib/python3.10/site-packages (from anyio<5.0,>=3.0->gradio) (1.2.0)
```

Code:-

```
import gradio as gr
from PIL import Image
import torch
import torch.nn as nn
from torchvision import models, transforms

# Load your trained model dynamically
def load_model(model_name, model_path):
    if model_name == "ResNet18":
        model = models.resnet18(pretrained=False) # Use
pretrained=False for your custom-trained models
        model.fc = nn.Linear(model.fc.in_features, 2) # Binary
classification
    elif model_name == "ResNet50":
        model = models.resnet50(pretrained=False)
        model.fc = nn.Linear(model.fc.in_features, 2) # Binary
classification
    elif model_name == "VGG16":
        model = models.vgg16(pretrained=False)
        model.classifier[6] =
nn.Linear(model.classifier[6].in_features, 2) # Binary classification
```



```

else:
    raise ValueError("Invalid model name.")

    # Load your trained model weights
    model.load_state_dict(torch.load(model_path,
map_location=torch.device('cpu'))))
    model.eval()
    return model

# Preprocessing function for input images
def preprocess_image(image):
    transform = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224,
0.225])
    ])
    return transform(image).unsqueeze(0) # Add batch dimension

# Prediction function
def predict(image, model_name):
    image = Image.fromarray(image) # Convert numpy array to PIL image
    input_tensor = preprocess_image(image)

    # Map model names to weight file paths
    model_paths = {
        "ResNet18":
"/kaggle/working/model/tumor_classification/ResNet18/ResNet18_final_epo
ch.pth",
        "ResNet50":
"/kaggle/working/model/tumor_classification/ResNet50/ResNet50_final_epo
ch.pth",
        "VGG16":
"/kaggle/working/model/tumor_classification/vgg16/vgg16_final_epoch.pth"
    }

    model_path = model_paths[model_name]
    model = load_model(model_name, model_path)

    with torch.no_grad():
        outputs = model(input_tensor)
        probabilities = torch.softmax(outputs, dim=1)[0]

    # Extract class names and their respective probabilities
    benign_prob = probabilities[0].item()
    malignant_prob = probabilities[1].item()
    predicted_idx = probabilities.argmax().item()

```

```

        predicted_class = class_names[predicted_idx]

    # Create a response string
    response = (
        f"Predicted Class: {predicted_class}\n"
        f"Benign Probability: {benign_prob:.2f}\n"
        f"Malignant Probability: {malignant_prob:.2f}"
    )
    return response

# Define global variables
class_names = ["benign", "malignant"] # Binary classes

# Define the Gradio interface
interface = gr.Interface(
    fn=predict,
    inputs=[
        gr.Image(type="numpy", label="Upload Image"),
        gr.Dropdown(choices=["ResNet18", "VGG16", "ResNet50"],
label="Select Model")
    ],
    outputs=gr.Textbox(label="Prediction"),
    title="Breast Cancer Classification",
    description="Upload an image and select a model to classify it as benign or malignant.",
)

# Launch the Gradio app
interface.launch()

```

<https://b489701ac1bf2e4c34.gradio.live>